

Cory Suzuki

STAT 574

Dr. Olga

12 April 2025

## Homework 4

### Problem 1

#### Python

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import yfinance as yf
from statistics import mean
import tensorflow
import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, GRU
from sklearn.metrics import mean_squared_error

# Historical Stock Extraction from Yahoo Finance

# data_download = yf.download('CVX', start='2010-01-01', end='2024-01-01')
# data_download = data_download.drop(columns=["High", "Low", "Open", "Volume"],
axis=1)
# print(data_download.head())

# data_download.to_csv('CVX_historical_data.csv')

chevron_data =
pd.read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/CVX_historical_
data.csv",
            index_col="Date", parse_dates=["Date"])

# Plotting daily Chevron stock closing prices

train = chevron_data[chevron_data.index < pd.to_datetime("2022-01-02",
format='%Y-%m-%d')]
test = chevron_data[chevron_data.index >= pd.to_datetime("2022-01-02",
format='%Y-%m-%d')]

plt.plot(train, color = "blue")
```

```

plt.plot(test, color = "green")
plt.ylabel('CVX Stock Price')
plt.xlabel('Date')
plt.title("Training and Testing Sets for CVX Stock Price Data")
plt.show()
# Rescaling Chevron data

chevron_data["Close_sc"] = (chevron_data["Close"]-
min(chevron_data["Close"]))/(max(chevron_data["Close"].values)-
min(chevron_data["Close"]))

train_set = chevron_data[chevron_data.index < pd.to_datetime("2022-01-02",
format='%Y-%m-%d')]
test_set = chevron_data[chevron_data.index >= pd.to_datetime("2022-01-02",
format='%Y-%m-%d')]
train_set = train_set.loc[:, "Close_sc"].values
test_set = test_set.loc[:, "Close_sc"].values

nsteps=60

def split_sequence(sequence, n_steps):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_ix = i + n_steps
        if end_ix > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_ix], sequence[end_ix]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y=split_sequence(train_set, nsteps)
# Fitting LSTM model

features = 1
train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

model_lstm = Sequential()
model_lstm.add(LSTM(units=6, activation="tanh", input_shape=(nsteps, features)))
model_lstm.add(Dense(units=1))

model_lstm.compile(loss="mse")
model_lstm.fit(train_x, train_y, epochs=5, batch_size=32)

```

```

# Creating testing set by adding nsteps observations from training set to testing
set
inputs=chevron_data.loc[:, "Close_sc"][len(chevron_data)-len(test_set)-
nsteps:].values
inputs=inputs.reshape(-1, 1)

# Splitting into samples
test_x, test_y=split_sequence(inputs, nsteps)

# Reshaping
test_x=test_x.reshape(test_x.shape[0], test_x.shape[1], features)

# Predicting for testing data
pred_y=model_lstm.predict(test_x)

# Inverse transforming the values
pred_y=pred_y*(max(chevron_data["Close"].values)-
min(chevron_data["Close"]))+min(chevron_data["Close"])
test_y=test_y*(max(chevron_data["Close"].values)-
min(chevron_data["Close"]))+min(chevron_data["Close"])

# Computing prediction accuracy
ind10=[]
ind15=[]
ind20=[]

for sub1, sub2 in zip(pred_y, test_y):
    ind10.append(1) if abs(sub1-sub2)<0.10*sub2 else ind10.append(0)
    ind15.append(1) if abs(sub1-sub2)<0.15*sub2 else ind15.append(0)
    ind20.append(1) if abs(sub1-sub2)<0.20*sub2 else ind20.append(0)

print('Accuracy within 10%:', round(mean(ind10),4))
print('Accuracy within 15%:', round(mean(ind15),4))
print('Accuracy within 20%:', round(mean(ind20),4))

#plotting actual and predicted values for testing data
plt.plot(test_y, color="green", label="actual price")
plt.plot(pred_y, color="orange", label="predicted price")
plt.title("Daily Chevron Stock Actual and Predicted Prices - LSTM Model")
plt.xlabel("Time")
plt.ylabel("Stock Price")
plt.legend()
plt.show()
# Fitting GRU Architecture Model

```

```

model_gru = Sequential()
model_gru.add(GRU(units=6, activation="tanh", input_shape=(nsteps, features)))
model_gru.add(Dense(units=1))

# Compiling the model
model_gru.compile(loss="mse")
model_gru.fit(train_x, train_y, epochs=5, batch_size=32)

# Predicting for testing data
pred_y=model_gru.predict(test_x)

# Inverse transforming the values
pred_y=pred_y*(max(chevron_data["Close"].values)-
min(chevron_data["Close"]))+min(chevron_data["Close"])

# Computing prediction accuracy
ind10=[]
ind15=[]
ind20=[]

for sub1, sub2 in zip(pred_y, test_y):
    ind10.append(1) if abs(sub1-sub2)<0.10*sub2 else ind10.append(0)
    ind15.append(1) if abs(sub1-sub2)<0.15*sub2 else ind15.append(0)
    ind20.append(1) if abs(sub1-sub2)<0.20*sub2 else ind20.append(0)

print('accuracy within 10%:', round(mean(ind10),4))
print('accuracy within 15%:', round(mean(ind15),4))
print('accuracy within 20%:', round(mean(ind20),4))

# Plotting actual and predicted values for testing data
plt.plot(test_y, color="green", label="actual price")
plt.plot(pred_y, color="orange", label="predicted price")
plt.title("Daily Chevron Stock Actual and Predicted Prices - GRU Model")
plt.xlabel("Time")
plt.ylabel("Stock Price")
plt.legend()
plt.show()

```

Training and Testing Sets for CVX Stock Price Data



Epoch 1/5

**93/93** ————— **4s** 15ms/step - loss: 0.0045

Epoch 2/5

**93/93** ————— **1s** 11ms/step - loss: 3.1114e-04

Epoch 3/5

**93/93** ————— **1s** 11ms/step - loss: 3.0604e-04

Epoch 4/5

**93/93** ————— **1s** 13ms/step - loss: 2.1896e-04

Epoch 5/5

**93/93** ————— **1s** 11ms/step - loss: 2.1180e-04


**24/24** ————— **0s** 13ms/step

Accuracy within 10%: 0.2709


Accuracy within 15%: 0.8194

Accuracy within 20%: 0.9987

Daily Chevron Stock Actual and Predicted Prices - LSTM Model



```
super().__init__(**kwargs)
93/93 ————— 4s 17ms/step - loss: 0.0524
Epoch 2/5
93/93 ————— 2s 19ms/step - loss: 0.0041
Epoch 3/5
93/93 ————— 2s 17ms/step - loss: 4.3160e-04
Epoch 4/5
93/93 ————— 2s 15ms/step - loss: 1.4043e-04
Epoch 5/5
93/93 ————— 2s 13ms/step - loss: 1.4531e-04
24/24 ————— 1s 16ms/step
accuracy within 10%: 0.9973
accuracy within 15%: 1
accuracy within 20%: 1
```



R

```
# STAT 574 HW4 Problem 1

# install.packages("keras3")
library(readr)
library(keras3)

cvx_data =
read.csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/CVX_historical_data.csv",
header=T, sep=",")

# Splitting data into training and testing sets.

cvx_data$Year = as.numeric(format(as.Date(cvx_data$Date, format="%Y-%m-%d"),
"%Y"))

train_data = cvx_data[which(cvx_data$Year<2022), 1:2]
```

```

test_data = cvx_data[which(cvx_data$Year>=2022), 1:2]

# Plotting training and testing data

plot(as.POSIXct(cvx_data$Date), cvx_data$Close, main="Daily Chevron Stock Closing
Prices",
xlab="Time", ylab="Stock Price", pch="", panel.first=grid())
lines(as.POSIXct(train_data$Date), train_data$Close, lwd=2, col="blue")
lines(as.POSIXct(test_data$Date), test_data$Close, lwd=2, col="green")
legend("topleft", c("training", "testing"), lty=1, lwd=2, col=c("blue", "green"))

# Scaling prices to fall within [0,1]

price = as.numeric(cvx_data$Close)
price_sc = (price-min(price))/(max(price)-min(price))

# Creating train_x and train_y

nsteps = 60
train_matrix = matrix(nrow=nrow(train_data)-nsteps, ncol=nsteps+1)
for (i in 1:(nrow(train_data)-nsteps)) {
  train_matrix[i,] = price_sc[i:(i+nsteps)]
}
train_x = array(train_matrix[,ncol(train_matrix)], dim=c(nrow(train_matrix),
nsteps, 1))
train_y = train_matrix[,ncol(train_matrix)]

# Creating test_x and test_y

test_matrix = matrix(nrow=nrow(test_data), ncol=nsteps+1)
for (i in 1:nrow(test_data)) {
  test_matrix[i,] =
price_sc[(i+nrow(train_matrix)):(i+nsteps+nrow(train_matrix))]
}
test_x = array(test_matrix[,ncol(test_matrix)], dim=c(nrow(test_matrix),
nsteps,1))
test_y = test_matrix[,ncol(test_matrix)]

# Fitting LSTM Model

LSTM_model = keras_model_sequential()

LSTM_model %>% layer_lstm(input_shape=dim(train_x)[2:3], units=nsteps)
LSTM_model %>% layer_dense(units=1, activation="tanh")
LSTM_model %>% compile(loss="mean_squared_error")

```



```

LSTM_model %>% fit(train_x, train_y, batch_size=32, epochs=5)
pred_y = LSTM_model %>% predict(test_x, batch_size=32)

test_y_re = test_y*(max(price)-min(price))+min(price)
pred_y_re = pred_y*(max(price)-min(price))+min(price)

# Computing prediction accuracy

accuracy10 = ifelse(abs(test_y_re-pred_y_re)<0.10*test_y_re, 1, 0)
accuracy15 = ifelse(abs(test_y_re-pred_y_re)<0.15*test_y_re, 1, 0)
accuracy20 = ifelse(abs(test_y_re-pred_y_re)<0.20*test_y_re, 1, 0)

print(paste("Accuracy within 10%", round(mean(accuracy10), 4)))
print(paste("Accuracy within 15%", round(mean(accuracy15), 4)))
print(paste("Accuracy within 20%", round(mean(accuracy20), 4)))

# Plotting actual and predicted values for testing data

plot(as.POSIXct(test_data$Date), test_y_re, type="l", lwd=2, col="green",
main="Daily Chevron Stock Actual and Predicted Prices - LSTM Model",
xlab="Time", ylab="Stock Price", panel.first=grid())
lines(as.POSIXct(test_data$Date), pred_y_re, lwd=2, col="orange")
legend("bottomright", c("actual", "predicted"), lty=1, lwd=2,
col=c("green", "orange"))

# Fitting GRU Model

gru_model = keras_model_sequential()

gru_model %>% layer_gru(input_shape=dim(train_x)[2:3], units=nsteps)
gru_model %>% layer_dense(units=1, activation="tanh")
gru_model %>% compile(loss="mean_squared_error")

gru_model %>% fit(train_x, train_y, batch_size=32, epochs=5)

pred_y_gru = gru_model %>% predict(test_x, batch_size=32)

pred_y_re_gru = pred_y_gru*(max(price)-min(price))+min(price)

# Computing prediction accuracy

accuracy10_gru = ifelse(abs(test_y_re-pred_y_re_gru)<0.10*test_y_re, 1, 0)
accuracy15_gru = ifelse(abs(test_y_re-pred_y_re_gru)<0.15*test_y_re, 1, 0)
accuracy20_gru = ifelse(abs(test_y_re-pred_y_re_gru)<0.20*test_y_re, 1, 0)


```

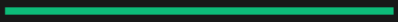
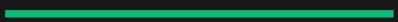
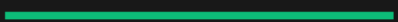



```
print(paste("Accuracy within 10%", round(mean(accuracy10_gru), 4)))
print(paste("Accuracy within 15%", round(mean(accuracy15_gru), 4)))
print(paste("Accuracy within 20%", round(mean(accuracy20_gru), 4)))

# Plotting actual and predicted values for testing data


plot(as.POSIXct(test_data$Date), test_y_re, type="l", lwd=2, col="green",
main="Daily Chevron Stock Actual and Predicted Prices - GRU Model",
xlab="Time", ylab="Stock Price", panel.first=grid())
lines(as.POSIXct(test_data$Date), pred_y_re_gru, lwd=2, col="orange")
legend("bottomright", c("actual", "predicted"), lty=1, lwd=2,
col=c("green", "orange"))
```

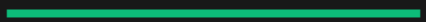




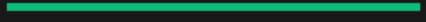
### Daily Chevron Stock CLOsing Prices




```
super().__init__(**kwargs)
Epoch 1/5
93/93  8s 48ms/step - loss: 0.0065
Epoch 2/5
93/93  6s 61ms/step - loss: 6.0162e-04
Epoch 3/5
93/93  6s 62ms/step - loss: 4.8289e-04
Epoch 4/5
93/93  6s 62ms/step - loss: 3.6237e-04
Epoch 5/5
93/93  6s 64ms/step - loss: 4.0895e-04
24/24  1s 39ms/step
[1] "Accuracy within 10%: 0.3347"
[1] "Accuracy within 15%: 0.9057"
[1] "Accuracy within 20%: 1"
```

### Daily Chevron Stock Actual and Predicted Prices - LSTM Model



```
Epoch 1/5
93/93  4s 23ms/step - loss: 0.0070
Epoch 2/5
93/93  3s 27ms/step - loss: 3.7110e-04
Epoch 3/5
93/93  3s 29ms/step - loss: 2.8554e-04
Epoch 4/5
93/93  2s 26ms/step - loss: 3.1382e-04
Epoch 5/5
93/93  3s 29ms/step - loss: 2.6224e-04
24/24  1s 19ms/step
[1] "Accuracy within 10%: 0.1633"
[1] "Accuracy within 15%: 0.8035"
[1] "Accuracy within 20%: 0.9987"
```



Problem 2

Python

```
import yfinance as yf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```

from sklearn.metrics import mean_squared_error
from statistics import mean
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, GRU

cvx_data =
pd.read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/CVX_historical_
data_shock.csv",
            index_col="Date", parse_dates=["Date"])
cvx_data = cvx_data.drop(["Close"], axis=1)

# Splitting the data into training and testing sets.

train = cvx_data[cvx_data.index < pd.to_datetime("2022-01-02", format='%Y-%m-
%d')]
test = cvx_data[cvx_data.index >= pd.to_datetime("2022-01-02", format='%Y-%m-
%d')]

train_set = train.loc[:, "Shock"].values
test_set = test.loc[:, "Shock"].values

# Splitting training data into samples.

nsteps=60

def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence)-1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)
# Fitting LSTM Model

features=1
train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

model_lstm = Sequential()
model_lstm.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))

```



```

model_lstm.add(Dense(units=1))

model_lstm.compile(loss="binary_crossentropy")
model_lstm.fit(train_x, train_y, epochs=5, batch_size=32)
inputs = cvx_data.loc[:, "Shock"][len(cvx_data.loc[:, "Shock"])-len(test_set)-
nsteps :].values

test_x, test_y = split_sequence(inputs)
test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)

pred_prob = model_lstm.predict(test_x)

cutoff = []
accuracy = []
for i in range(99):
    tp=0
    tn=0
    cutoff.append(0.01*(i+1))
    for sub1, sub2 in zip(pred_prob, test_y):
        tp_ind = 1 if (sub1>0.01*(i+1) and sub2==1) else 0
        tn_ind = 1 if (sub1<0.01*(i+1) and sub2==0) else 0
        tp+=tp_ind
        tn+=tn_ind
    accuracy_i = (tp+tn)/len(pred_prob)
    accuracy.append(accuracy_i)

df = pd.DataFrame({'accuracy': accuracy, 'cut-off':cutoff})
max_accuracy = max(accuracy)
optimal=df[df['accuracy']==max_accuracy]
print(optimal)
# Fitting GRU Architecture

model_gru = Sequential()
model_gru.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps, features)))
model_gru.add(Dense(units=1))

model_gru.compile(loss="binary_crossentropy")
model_gru.fit(train_x, train_y, epochs=5, batch_size=32)

pred_prob = model_gru.predict(test_x)

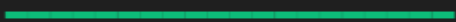
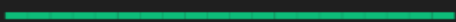
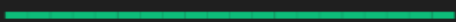
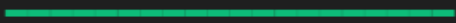
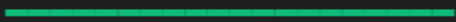
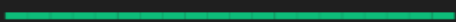
cutoff=[]
accuracy=[]
for i in range(99):
    tp=0

```

```
tn=0
cutoff.append(0.01*(i+1))
for sub1, sub2 in zip(pred_prob, test_y):
    tp_ind=1 if (sub1>0.01*(i+1) and sub2==1) else 0
    tn_ind=1 if (sub1<0.01*(i+1) and sub2==0) else 0
    tp+=tp_ind
    tn+=tn_ind

accuracy_i=(tp+tn)/len(pred_prob)
accuracy.append(accuracy_i)

df=pd.DataFrame({'accuracy': accuracy, 'cut-off': cutoff})
max_accuracy=max(accuracy)
optimal=df[df['accuracy']==max_accuracy]
print(optimal)
```

```
super().__init__(**kwargs)
93/93  2s 9ms/step - loss: 0.5266
Epoch 2/5
93/93  1s 9ms/step - loss: 0.4353
Epoch 3/5
93/93  1s 9ms/step - loss: 0.4260
Epoch 4/5
93/93  1s 9ms/step - loss: 0.4176
Epoch 5/5
93/93  1s 9ms/step - loss: 0.3981
16/16  0s 13ms/step
```

	accuracy	cut-off
0	0.954092	0.01
1	0.954092	0.02
2	0.954092	0.03
3	0.954092	0.04
4	0.954092	0.05
..	...	...
74	0.954092	0.75
75	0.954092	0.76
76	0.954092	0.77
77	0.954092	0.78
78	0.954092	0.79

```

super().__init__(**kwargs)
93/93 ————— 3s 16ms/step - loss: 2.4998
Epoch 2/5
93/93 ————— 1s 16ms/step - loss: 0.6285
Epoch 3/5
93/93 ————— 2s 16ms/step - loss: 0.4338
Epoch 4/5
93/93 ————— 2s 18ms/step - loss: 0.4110
Epoch 5/5
93/93 ————— 2s 18ms/step - loss: 0.4252
16/16 ————— 1s 19ms/step

```

	accuracy	cut-off
0	0.954092	0.01
1	0.954092	0.02
2	0.954092	0.03
3	0.954092	0.04
4	0.954092	0.05
..	...	...
80	0.954092	0.81
81	0.954092	0.82
82	0.954092	0.83
83	0.954092	0.84
84	0.954092	0.85

R

```

# STAT 574 HW4 Problem 2

library(readr)
library(dplyr)
library(keras3)

cvx_data =
read.csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/CVX_historical_data_shock.csv",
header=T, sep=",")

```

```

# Splitting data into training and testing sets.

cvx_data$Year = as.numeric(format(as.Date(cvx_data$Date, format="%m/%d/%Y"),
"%Y"))
train_data = cvx_data[which(cvx_data$Year<2022), 1:2]
test_data = cvx_data[which(cvx_data$Year>=2022), 1:2]

nsteps = 60
train_matrix = matrix(nrow=nrow(train_data)-nsteps, ncol=nsteps+1)
for (i in 1:(nrow(train_data)-nsteps)) {
  train_matrix[i,] = cvx_data$Shock[i:(i+nsteps)]
}
train_x = array(train_matrix[,ncol(train_matrix)], dim=c(nrow(train_matrix),
nsteps, 1))
train_y = train_matrix[,ncol(train_matrix)]

# Creating test_x and test_y

test_matrix = matrix(nrow=nrow(test_data), ncol=nsteps+1)
for (i in 1:nrow(test_data)) {
  test_matrix[i,] =
cvx_data$Shock[(i+nrow(train_matrix)):(i+nsteps+nrow(train_matrix))]
}
test_x = array(test_matrix[,ncol(test_matrix)], dim=c(nrow(test_matrix), nsteps,
1))
test_y = test_matrix[,ncol(test_matrix)]

# Fitting LSTM Model

LSTM_biclass = keras_model_sequential()
LSTM_biclass %>% layer_dense(input_shape=dim(train_x)[2:3], units=nsteps)
LSTM_biclass %>% layer_lstm(units=25)
LSTM_biclass %>% layer_dense(units=1, activation="sigmoid")
LSTM_biclass %>% compile(loss="binary_crossentropy")

LSTM_biclass %>% fit(train_x, train_y, batch_size=32, epochs=5)

# Computing prediction accuracy for testing data.

pred_prob = LSTM_biclass %>% predict(test_x)
match = cbind(test_y, pred_prob)
tp = matrix(NA, nrow=nrow(match), ncol=99)
tn = matrix(NA, nrow=nrow(match), ncol=99)

for (i in 1:99) {

```

```

    tp[,i] = ifelse(match[,1]==1 & match[,2]>0.01*i, 1, 0)
    tn[,i] = ifelse(match[,1]==0 & match[,2]<=0.01*i, 1, 0)
}

trueclassrate = matrix(NA, nrow=99, ncol=2)
for (i in 1:99) {
    trueclassrate[i, 1] = 0.01*i
    trueclassrate[i, 2] = sum(tp[,i]+tn[,i])/nrow(match)
}

print(trueclassrate[which(trueclassrate[,2]==max(trueclassrate[,2])),])

# Fitting GRU model

gru_biclass = keras_model_sequential()
gru_biclass %>% layer_dense(input_shape=dim(train_x)[2:3], units=nsteps)
gru_biclass %>% layer_gru(units=25)
gru_biclass %>% layer_dense(units=1, activation="sigmoid")
gru_biclass %>% compile(loss="binary_crossentropy")

gru_biclass %>% fit(train_x, train_y, batch_size=32, epochs=5)

# Computing prediction accuracy for testing data.

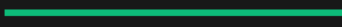
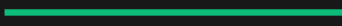
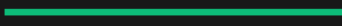

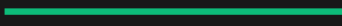

pred_prob_gru = gru_biclass %>% predict(test_x)
match_gru = cbind(test_y, pred_prob_gru)
tp_gru = matrix(NA, nrow=nrow(match_gru), ncol=99)
tn_gru = matrix(NA, nrow=nrow(match_gru), ncol=99)

for (i in 1:99) {
    tp[,i] = ifelse(match_gru[,1]==1 & match_gru[,2]>0.01*i, 1, 0)
    tn[,i] = ifelse(match_gru[,1]==0 & match_gru[,2]<=0.01*i, 1, 0)
}

trueclassrate_gru = matrix(NA, nrow=99, ncol=2)
for (i in 1:99) {
    trueclassrate_gru[i, 1] = 0.01*i
    trueclassrate_gru[i, 2] = sum(tp[,i]+tn[,i])/nrow(match_gru)
}

print(trueclassrate_gru[which(trueclassrate[,2]==max(trueclassrate[,2])),])

```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/5
93/93  9s 53ms/step - loss: 0.4590
Epoch 2/5
93/93  5s 52ms/step - loss: 0.4235
Epoch 3/5
93/93  5s 56ms/step - loss: 0.4255
Epoch 4/5
93/93  5s 52ms/step - loss: 0.4224
Epoch 5/5
93/93  5s 54ms/step - loss: 0.4137
16/16  2s 66ms/step
```

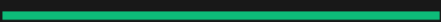
	[,1]	[,2]
[1,]	0.01	0.9540918
[2,]	0.02	0.9540918
[3,]	0.03	0.9540918
[4,]	0.04	0.9540918
[5,]	0.05	0.9540918
[6,]	0.06	0.9540918
[7,]	0.07	0.9540918
[8,]	0.08	0.9540918
[9,]	0.09	0.9540918
[10,]	0.10	0.9540918
[11,]	0.11	0.9540918
[12,]	0.12	0.9540918
[13,]	0.13	0.9540918
[14,]	0.14	0.9540918
[15,]	0.15	0.9540918
[16,]	0.16	0.9540918
[17,]	0.17	0.9540918
[18,]	0.18	0.9540918
[19,]	0.19	0.9540918
[20,]	0.20	0.9540918
[21,]	0.21	0.9540918



```
[22,] 0.22 0.9540918
[23,] 0.23 0.9540918
[24,] 0.24 0.9540918
[25,] 0.25 0.9540918
[26,] 0.26 0.9540918
[27,] 0.27 0.9540918
[28,] 0.28 0.9540918
[29,] 0.29 0.9540918
[30,] 0.30 0.9540918
[31,] 0.31 0.9540918
[32,] 0.32 0.9540918
[33,] 0.33 0.9540918
[34,] 0.34 0.9540918
[35,] 0.35 0.9540918
[36,] 0.36 0.9540918
[37,] 0.37 0.9540918
[38,] 0.38 0.9540918
[39,] 0.39 0.9540918
[40,] 0.40 0.9540918
[41,] 0.41 0.9540918
[42,] 0.42 0.9540918
[43,] 0.43 0.9540918
[44,] 0.44 0.9540918
[45,] 0.45 0.9540918
[46,] 0.46 0.9540918
[47,] 0.47 0.9540918
[48,] 0.48 0.9540918
[49,] 0.49 0.9540918
[50,] 0.50 0.9540918
```

```
[51,] 0.51 0.9540918
[52,] 0.52 0.9540918
[53,] 0.53 0.9540918
[54,] 0.54 0.9540918
[55,] 0.55 0.9540918
[56,] 0.56 0.9540918
[57,] 0.57 0.9540918
[58,] 0.58 0.9540918
[59,] 0.59 0.9540918
[60,] 0.60 0.9540918
[61,] 0.61 0.9540918
[62,] 0.62 0.9540918
[63,] 0.63 0.9540918
[64,] 0.64 0.9540918
[65,] 0.65 0.9540918
[66,] 0.66 0.9540918
[67,] 0.67 0.9540918
[68,] 0.68 0.9540918
[69,] 0.69 0.9540918
[70,] 0.70 0.9540918
[71,] 0.71 0.9540918
[72,] 0.72 0.9540918
[73,] 0.73 0.9540918
[74,] 0.74 0.9540918
[75,] 0.75 0.9540918
[76,] 0.76 0.9540918
[77,] 0.77 0.9540918
[78,] 0.78 0.9540918
[79,] 0.79 0.9540918
```

Epoch 1/5

**93/93**  **5s** 23ms/step - loss: 0.4926

Epoch 2/5

**93/93**  **2s** 25ms/step - loss: 0.4268

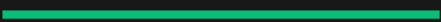
Epoch 3/5

**93/93**  **2s** 22ms/step - loss: 0.4262

Epoch 4/5

**93/93**  **2s** 22ms/step - loss: 0.4215

Epoch 5/5

**93/93**  **2s** 23ms/step - loss: 0.4116

**16/16**  **1s** 27ms/step

	[,1]	[,2]
[1,]	0.01	0.9540918
[2,]	0.02	0.9540918
[3,]	0.03	0.9540918
[4,]	0.04	0.9540918
[5,]	0.05	0.9540918
[6,]	0.06	0.9540918
[7,]	0.07	0.9540918
[8,]	0.08	0.9540918
[9,]	0.09	0.9540918
[10,]	0.10	0.9540918
[11,]	0.11	0.9540918
[12,]	0.12	0.9540918
[13,]	0.13	0.9540918
[14,]	0.14	0.9540918
[15,]	0.15	0.9540918
[16,]	0.16	0.9540918
[17,]	0.17	0.9540918
[18,]	0.18	0.9540918
[19,]	0.19	0.9540918
[20,]	0.20	0.9540918
[21,]	0.21	0.9540918
[22,]	0.22	0.9540918

```
[23,] 0.23 0.9540918
[24,] 0.24 0.9540918
[25,] 0.25 0.9540918
[26,] 0.26 0.9540918
[27,] 0.27 0.9540918
[28,] 0.28 0.9540918
[29,] 0.29 0.9540918
[30,] 0.30 0.9540918
[31,] 0.31 0.9540918
[32,] 0.32 0.9540918
[33,] 0.33 0.9540918
[34,] 0.34 0.9540918
[35,] 0.35 0.9540918
[36,] 0.36 0.9540918
[37,] 0.37 0.9540918
[38,] 0.38 0.9540918
[39,] 0.39 0.9540918
[40,] 0.40 0.9540918
[41,] 0.41 0.9540918
[42,] 0.42 0.9540918
[43,] 0.43 0.9540918
[44,] 0.44 0.9540918
[45,] 0.45 0.9540918
[46,] 0.46 0.9540918
[47,] 0.47 0.9540918
[48,] 0.48 0.9540918
[49,] 0.49 0.9540918
[50,] 0.50 0.9540918
[51,] 0.51 0.9540918
```

```
[52,] 0.52 0.9540918
[53,] 0.53 0.9540918
[54,] 0.54 0.9540918
[55,] 0.55 0.9540918
[56,] 0.56 0.9540918
[57,] 0.57 0.9540918
[58,] 0.58 0.9540918
[59,] 0.59 0.9540918
[60,] 0.60 0.9540918
[61,] 0.61 0.9540918
[62,] 0.62 0.9540918
[63,] 0.63 0.9540918
[64,] 0.64 0.9540918
[65,] 0.65 0.9540918
[66,] 0.66 0.9540918
[67,] 0.67 0.9540918
[68,] 0.68 0.9540918
[69,] 0.69 0.9540918
[70,] 0.70 0.9540918
[71,] 0.71 0.9540918
[72,] 0.72 0.9540918
[73,] 0.73 0.9540918
[74,] 0.74 0.9540918
[75,] 0.75 0.9540918
[76,] 0.76 0.9540918
[77,] 0.77 0.9540918
[78,] 0.78 0.9540918
[79,] 0.79 0.9540918
```

Problem 3

## Python

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statistics import mean
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM, GRU

seattle_weather =
pd.read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/w
eather_description.csv",
index_col = "datetime", parse_dates=["datetime"])
seattle_weather['Seattle'].value_counts()
code_conditions = {"sky is clear": "clear", "light rain": "rain", "overcast
clouds": "cloudy",
                  "broken clouds": "cloudy", "mist": "fog", "scattered clouds":
"cloudy",
                  "few clouds": "cloudy", "moderate rain": "rain", "light
intensity drizzle": "rain",
                  "fog": "fog", "haze": "fog", "heavy intensity rain": "cloudy",
"smoke": "cloudy",
                  "light snow": "snow", "light intensity shower rain": "rain",
                  "proximity thunderstorm": "rain", "very heavy rain": "rain",
"drizzle": "rain",
                  "thunderstorm": "rain", "thunderstorm with light rain":
"rain",
                  "heavy intensity drizzle": "rain", "heavy snow": "snow",
"shower rain": "rain",
                  "thunderstorm with heavy rain": "rain", "thunderstorm with
rain": "rain",
                  "light shower snow": "snow", "squalls": "rain", "heavy
intensity shower rain": "rain"}

seattle_weather['Seattle'] = seattle_weather['Seattle'].map(code_conditions)
seattle_weather = pd.get_dummies(seattle_weather['Seattle'])
# LSTM Models

# Rain

time_start = 2012
time_end = 2016
def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "rain"].values
    test = seattle_weather.loc[f"{time_end+1}":, "rain"].values
```

```

        return train, test

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

features = 1
train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_rain = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_rain = fitted_model.predict(test_x)
# Fog Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "fog"].values
    test = seattle_weather.loc[f"{time_end+1}":, "fog"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

def split_sequence(sequence):
    x, y= list(), list()
    for i in range(len(sequence)):

```



```

        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_y = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_fog = fitted_model.predict(test_x)
# Cloudy Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "cloudy"].values
    test = seattle_weather.loc[f"{time_end+1}":, "cloudy"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

```

```

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_cloudy = split_sequence(test_set)
test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)

pred_prob_cloudy = fitted_model.predict(test_x)
# Snow Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "snow"].values
    test = seattle_weather.loc[f"{time_end+1}":, "snow"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")

```

```

fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)
test_x, test_snow = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_snow = fitted_model.predict(test_x)
# Clear Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "clear"].values
    test = seattle_weather.loc[f"{time_end+1}":, "clear"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(LSTM(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)
test_x, test_clear = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_clear = fitted_model.predict(test_x)
# Computing prediction accuracy for LSTM models

pred_prob_all = np.concatenate((pred_prob_rain, pred_prob_fog, pred_prob_clear,
pred_prob_cloudy, pred_prob_snow), axis=1)
pred_prob_all = pd.DataFrame(pred_prob_all)

```

```

pred_class = pred_prob_all.idxmax(axis=1)

test_all = np.c_[test_rain, test_fog, test_clear, test_cloudy, test_snow]
test_all = pd.DataFrame(test_all)
true_class = test_all.idxmax(axis=1)

match = []
for i in range(len(pred_class)):
    if pred_class[i] == true_class[i]:
        match.append(1)
    else:
        match.append(0)

print(round(mean(match), 4))

```

0.7094

```

# GRU Models

# Rain Model

time_start = 2012
time_end = 2016

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "rain"].values
    test = seattle_weather.loc[f"{time_end+1}":, "rain"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

```

```

features = 1
train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_rain = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_rain = fitted_model.predict(test_x)
# Fog Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "fog"].values
    test = seattle_weather.loc[f"{time_end+1}":, "fog"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

```

```

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_fog = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_fog = fitted_model.predict(test_x)
# Cloudy Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "cloudy"].values
    test = seattle_weather.loc[f"{time_end+1}":, "cloudy"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_fog = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)

```

```

pred_prob_fog = fitted_model.predict(test_x)
# Snow Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "snow"].values
    test = seattle_weather.loc[f"{time_end+1}":, "snow"].values
    return train, test

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_fog = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_fog = fitted_model.predict(test_x)
# Clear Model

def train_test_split(time_start, time_end):
    train = seattle_weather.loc[f"{time_start}":f"{time_end}", "clear"].values
    test = seattle_weather.loc[f"{time_end+1}":, "clear"].values
    return train, test

```

```

train_set, test_set = train_test_split(time_start, time_end)

nsteps = 60
def split_sequence(sequence):
    x, y = list(), list()
    for i in range(len(sequence)):
        end_i = i + nsteps
        if end_i > len(sequence) - 1:
            break
        seq_x, seq_y = sequence[i:end_i], sequence[end_i]
        x.append(seq_x)
        y.append(seq_y)
    return np.array(x), np.array(y)

train_x, train_y = split_sequence(train_set)

train_x = train_x.reshape(train_x.shape[0], train_x.shape[1], features)

fitted_model = Sequential()
fitted_model.add(GRU(units=6, activation="sigmoid", input_shape=(nsteps,
features)))
fitted_model.add(Dense(units=1, activation="sigmoid"))

fitted_model.compile(loss="binary_crossentropy")
fitted_model.fit(train_x, train_y, epochs=5, batch_size=32)

test_x, test_fog = split_sequence(test_set)

test_x = test_x.reshape(test_x.shape[0], test_x.shape[1], features)
pred_prob_fog = fitted_model.predict(test_x)
# Computing prediction accuracy for GRU model.

pred_prob_all = np.concatenate((pred_prob_rain, pred_prob_fog, pred_prob_cloudy,
pred_prob_snow, pred_prob_clear), axis=1)
pred_prob_all = pd.DataFrame(pred_prob_all)
pred_class = pred_prob_all.idxmax(axis=1)

test_all = np.c_[test_rain, test_fog, test_cloudy, test_snow, test_clear]
test_all = pd.DataFrame(test_all)
true_class = test_all.idxmax(axis=1)

match_gru = []
for i in range(len(pred_class)):
    if pred_class[i] == true_class[i]:
        match_gru.append(1)

```



```

    else:
        match_gru.append(0)

print(round(mean(match_gru), 4))

```

0.6338

R

```

# STAT 574 HW4 Problem 3

library(readr)
library(keras3)

weather_data =
read.csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/weat
her_description.csv")
table(DT <- weather_data$Seattle)

DT = ifelse(DT=="sky is clear", "clear", ifelse(DT %in% c("broken clouds", "few
clouds", "overcast clouds", "scattered clouds", "smoke"), "cloudy",
ifelse(DT %in% c("heavy shower snow", "heavy snow", "light shower snow", "light
snow", "shower snow", "snow"), "snow",
ifelse(DT %in% c("fog", "haze", "mist"), "fog", "rain"))))

weather_data$clear = ifelse(DT=="clear", 1, 0)
weather_data$cloudy = ifelse(DT=="cloudy",1,0)
weather_data$rain = ifelse(DT=="rain",1,0)
weather_data$snow = ifelse(DT=="snow",1,0)
weather_data$fog = ifelse(DT=="fog",1,0)
weather_data$year = format(as.Date(weather_data$datetime, format="%m/%d/%y"),
"%Y")

rnn_model = function(modelname, varname) {
  train_data = weather_data[which(weather_data$year<2017), varname]
  test_data = weather_data[which(weather_data$year==2017), varname]
  nsteps = 60
  train_matrix = matrix(nrow=length(train_data)-nsteps, ncol=nsteps+1)
  for (i in 1:(length(train_data)-nsteps)) {
    train_matrix[i,] = weather_data[i:(i+nsteps), varname]
  }
  train_x = array(train_matrix[,-ncol(train_matrix)], dim=c(nrow(train_matrix),
nsteps, 1))

```

```

train_y = train_matrix[,ncol(train_matrix)]
test_matrix = matrix(nrow=length(test_data), ncol=nsteps+1)
for (i in 1:length(test_data)) {
  test_matrix[i,] =
weather_data[(i+nrow(train_matrix)):(i+nsteps+nrow(train_matrix)), varname]
}
test_x = array(test_matrix[,ncol(test_matrix)], dim=c(nrow(test_matrix),
nsteps, 1))
test_y = test_matrix[,ncol(test_matrix)]

fitted_model = keras_model_sequential()
fitted_model %>% layer_dense(input_shape=dim(train_x)[2:3], units=nsteps)
if (modelname=="lstm") {
  fitted_model %>% layer_lstm(units=6)
} else {
  fitted_model %>% layer_gru(units=6)
}
fitted_model %>% layer_dense(units=1, activation="sigmoid")
fitted_model %>% compile(loss='binary_crossentropy')

fitted_model %>% fit(train_x, train_y, batch_size=32, epochs=1)
pred_prob = fitted_model %>% predict(test_x)
return(list(test_y, pred_prob))
}

accuracy = function() {
  test_y = bind_cols(test_clear, test_cloudy, test_snow, test_fog, test_rain)
  colnames(test_y) = 1:5
  true_class = as.numeric(apply(test_y, 1, function(x)
  colnames(test_y)[which.max(x)]))
  pred_prob = bind_cols(pred_prob_clear, pred_prob_cloudy, pred_prob_snow,
pred_prob_fog, pred_prob_rain)
  colnames(pred_prob) = 1:5
  pred_class = as.numeric(apply(pred_prob, 1, function(x)
  colnames(pred_prob)[which.max(x)]))
  match = c()
  for (i in 1:length(pred_class)) {
    match[i] = ifelse(pred_class[i] == true_class[i],1,0)
  }
  return(round(mean(match), 4))
}

# Running LSTM Binary Classification Models

list_clear = rnn_model('lstm', 'clear')

```

```
test_clear = list_clear[1]
pred_prob_clear = list_clear[2]

list_cloudy = rnn_model('lstm', 'cloudy')
test_cloudy = list_cloudy[1]
pred_prob_cloudy = list_cloudy[2]

list_snow = rnn_model('lstm', 'snow')
test_snow = list_snow[1]
pred_prob_snow = list_snow[2]

list_fog = rnn_model('lstm', 'fog')
test_fog = list_fog[1]
pred_prob_fog = list_fog[2]

list_rain = rnn_model('lstm', 'rain')
test_rain = list_rain[1]
pred_prob_rain = list_rain[2]
```

```
## [1] 0.71638
```

```
# Running GRU Binary Classification Models
```

```
list_clear = rnn_model('gru', 'clear')
test_clear = list_clear[1]
pred_prob_clear = list_clear[2]

list_cloudy = rnn_model('gru', 'cloudy')
test_cloudy = list_cloudy[1]
pred_prob_cloudy = list_cloudy[2]

list_snow = rnn_model('gru', 'snow')
test_snow = list_snow[1]
pred_prob_snow = list_snow[2]

list_fog = rnn_model('gru', 'fog')
test_fog = list_fog[1]
pred_prob_fog = list_fog[2]

list_rain = rnn_model('gru', 'rain')
test_rain = list_rain[1]
pred_prob_rain = list_rain[2]
```

```
## [1] 0.7105
```

#### Problem 4

R

```
# STAT 574 HW4 Problem 4

install.packages("changepoint")
library(readr)
library(changepoint)

gold_data =
read.csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/wheat_data.csv",
header=T, sep=",")

# Detection of change points for change in mean.

mean_det = cpt.mean(gold_data$Close, penalty="AIC", method="BinSeg", Q=3)
plot(mean_det, cpt.col="red", ylab="Daily Closing Price", main="Change Point
Detection for Change in Mean")
paste("Change Point Locations: ", paste(mean_det@cpts, collapse=", "))

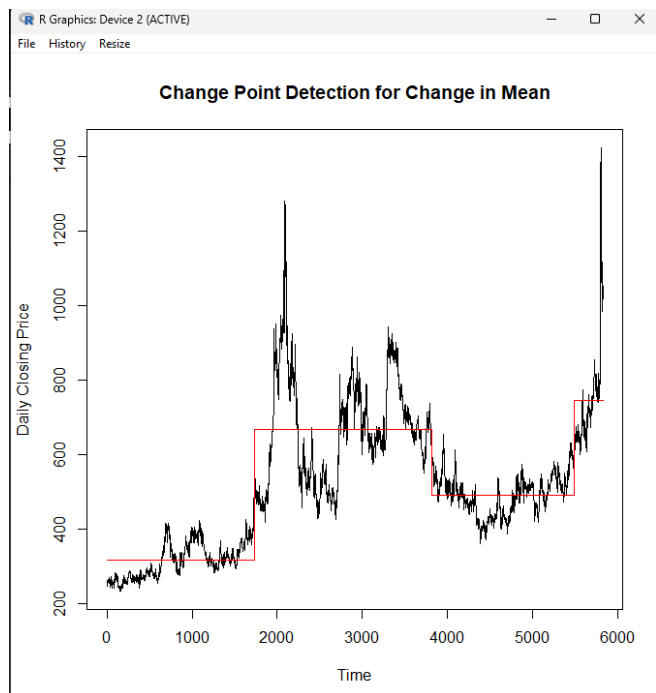
# Detection of change points for change in variance.

var_det = cpt.var(gold_data$Close, penalty="AIC", method="BinSeg", Q=3)
plot(var_det, cpt.col="red", ylab="Daily Closing Price", main="Change Point
Detection for Change in Variance")
paste("Change Point Locations: ", paste(var_det@cpts, collapse=", "))


# Detection of Change Points for change in mean and variance.

mean_var_det = cpt.meanvar(gold_data$Close, penalty="AIC", method="BinSeg", Q=3)
plot(mean_var_det, cpt.col="red", ylab="Daily Closing Price",
main="Change Point Detection for Change in Mean and Variance")
paste("Change Point Locations: ", paste(mean_var_det@cpts, collapse=", "))
```


Change in Mean



## Change in Variance



## Change in Mean and Variance



## Problem 5

R

```
# install.packages("tibbletime")
# install.packages("anomalize")
# install.packages("tidyverse")
library(readr)
library(tibbletime)
library(anomalize)
library(tidyverse)

wheat_data =
  read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/wheat_data.csv",
    header=T, sep=",")


wheat_data$Date = as.Date(wheat_data$Date, "%m/%d/%Y")

wheat_data_tbl = as_tbl_time(wheat_data, Date)

print(wheat_data_tbl %>% time_decompose(Close, method="stl") %>%
  anomalize(remainder,
    method="iqr") %>% time_recompose() %>% plot_anomalies(time_recomposed=T,
    color_no='blue',
    color_yes='red', fill_ribbon='skyblue', size_circles=4) +
  labs(title="Anomalies in Daily Closing Prices of Wheat", subtitle="1/4/2000-4/8/2022"))
```

## Anomalies in Daily Closing Prices of Wheat

1/4/2000-4/8/2022



## Problem 6

R

```
# install.packages("BiocManager")  
# BiocManager::install("EBImage")  
library(keras3)  
library(EBImage)
```

```
# Preparing Data

# Caracals

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/CARACALS")
img_caracals = sample(dir());
train_caracals = list(NULL);
for (i in 1:length(img_caracals)) {
  train_caracals[[i]] = readImage(img_caracals[i])
  train_caracals[[i]] = resize(train_caracals[[i]], 100, 100)
}

# Cheetahs

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/CHEETAHS")
img_cheetahs = sample(dir());
train_cheetahs = list(NULL);
for (i in 1:length(img_cheetahs)) {
  train_cheetahs[[i]] = readImage(img_cheetahs[i])
  train_cheetahs[[i]] = resize(train_cheetahs[[i]], 100, 100)
}

# Lions

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/LIONS")
img_lions = sample(dir());
train_lions = list(NULL);
for (i in 1:length(img_lions)) {
  train_lions[[i]] = readImage(img_lions[i])
  train_lions[[i]] = resize(train_lions[[i]], 100, 100)
}

# Tigers

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/TIGERS")
img_tigers = sample(dir());
train_tigers = list(NULL);
for (i in 1:length(img_tigers)) {
  train_tigers[[i]] = readImage(img_tigers[i])
  train_tigers[[i]] = resize(train_tigers[[i]], 100, 100)
}
```



```

train_pool = c(train_caracals[1:40], train_cheetahs[1:40], train_lions[1:40],
train_tigers[1:40])

# Permutation of image dimensions.

train = aperm(combine(train_pool), c(4,1,2,3))

# Creating image labels.

train_y = c(rep(0,40), rep(1,40), rep(2,40), rep(3,40))
train_lab = to_categorical(train_y)

# Preparing testing data.

# Caracals

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/CARACALS")
img_caracals = sample(dir())
test_caracals = list(NULL)

for (i in 1:length(img_caracals)) {
  test_caracals[[i]] = readImage(img_caracals[i])
  test_caracals[[i]] = resize(test_caracals[[i]], 100, 100)
}

# Cheetahs

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/CHEETAHS")
img_cheetahs = sample(dir())
test_cheetahs = list(NULL)

for (i in 1:length(img_cheetahs)) {
  test_cheetahs[[i]] = readImage(img_cheetahs[i])
  test_cheetahs[[i]] = resize(test_cheetahs[[i]], 100, 100)
}

# Lions

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/LIONS")
img_lions = sample(dir())
test_lions = list(NULL)

```

```

for (i in 1:length(img_lions)) {
  test_lions[[i]] = readImage(img_lions[i])
  test_lions[[i]] = resize(test_lions[[i]], 100, 100)
}

# Tigers

setwd("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/WildAnimalsImages/train/TIGERS")
img_tigers = sample(dir())
test_tigers = list(NULL)

for (i in 1:length(img_tigers)) {
  test_tigers[[i]] = readImage(img_tigers[i])
  test_tigers[[i]] = resize(test_tigers[[i]], 100, 100)
}

test_pool = c(test_caracals[1:3], test_cheetahs[1:3], test_lions[1:3],
test_tigers[1:3])
test = aperm(combine(test_pool), c(4,1,2,3))
test_y = c(rep(0,3), rep(1,3), rep(2,3), rep(3,3))
test_lab = to_categorical(test_y)

# Fitting CNN Architecture Model.

model_cnn = keras_model_sequential()

model_cnn %>% layer_conv_2d(filters=40, kernel_size=c(3,3),
activation="relu", input_shape=c(100,100,3)) %>%
layer_conv_2d(filters=40, kernel_size=c(3,3), activation="relu") %>%
layer_max_pooling_2d(pool_size=c(3,3)) %>% layer_dropout(rate=0.25) %>%
layer_conv_2d(filters=80, kernel_size=c(3,3), activation="relu") %>%
layer_conv_2d(filters=80, kernel_size=c(3,3), activation="relu") %>%
layer_max_pooling_2d(pool_size=c(3,3)) %>% layer_dropout(rate=0.35) %>%
layer_flatten() %>% layer_dense(units=256, activation="relu") %>%
layer_dropout(rate=0.25) %>% layer_dense(units=4, activation="softmax") %>%
compile(loss="categorical_crossentropy", optimizer=optimizer_adam(),
metrics=c("accuracy"))

history = model_cnn %>% fit(train, train_lab, epochs=50, batch_size=40,
validation_split=0.2)

# Computing prediction accuracy for testing set.

```

```

model_cnn %>% evaluate(test, test_lab)

pred_class = as.array(model_cnn %>% predict(test) %>% k_argmax())
print(pred_class)
print(test_y)

print(paste("accuracy= ", round(1-mean(test_y!=pred_class), digits=4)))

```

```
## [1] "accuracy= 0.9167"
```

Problem 7

R

```

library(readr)
library(dplyr)
library(gutenbergr)
library(stringr)
library(tidytext)
library(stopwords)
library(tibble)
library(wordcloud)
library(ggplot2)

# Book selected: The White Company by Sir Arthur Conan Doyle

book_sel = gutenbergr_download(903, meta_fields="author")

book_sel = as_tibble(book_sel %>%
mutate(document=row_number()) %>%
select(-gutenberg_id))

tidy_book = book_sel %>% unnest_tokens(word, text) %>%
group_by(word) %>% filter(n()>10) %>%
ungroup()

# Identifying and dropping stopwords from text.

stopword = as_tibble(stopwords::stopwords("en"))
stopword = rename(stopword, word=value)
tb = anti_join(tidy_book, stopword, by='word')

```

```
word_count = tb %>% count(word, sort=T)
print(word_count)

# Displaying top 25 words.

tb %>%
count(author, word, sort=T) %>%
filter(n>=190) %>%
mutate(word=reorder(word,n)) %>%
ggplot(aes(word,n)) +
geom_col(aes(fill=author)) +
xlab(NULL) +
scale_y_continuous(expand=c(0, 0)) +
coord_flip() +
theme_classic(base_size=12) +
labs(fill="Author", title="Word Frequency", subtitle="25 top words")+
theme(plot.title = element_text(lineheight=.8, face="bold"))+
scale_fill_brewer()


# Plotting word cloud.

tb %>%
count(word) %>%
with(wordcloud(word, n, max.words=25, colors=brewer.pal(10, "Set1")))
```

```
# A tibble: 25 × 2
  word      n
  <chr>   <int>
1 upon    931
2 sir     719
3 said    629
4 man     514
5 alleyne 480
6 one     458
7 cried   398
8 nigel   350
9 may     348
10 two    325
# i 15 more rows
```

## Word Frequency

25 top words



upon  
sir  
said  
back  
john  
may  
eyes  
well  
US  
good  
hath  
man  
one  
see  
nigel  
great  
long  
now  
yet  
two  
cried  
men  
came  
alleyne

---

## Problem 8

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
```

```

from wordcloud import WordCloud
from simpletransformers.classification import ClassificationModel

air_data =
pd.read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/U
SAirlinesTweets.csv")
air_data.drop_duplicates(subset=['tweet'], keep='first', inplace=True)
text = " ".join([x for x in air_data.tweet[air_data.sentiment == 'positive']])


# Positive Tweets

# Plotting wordclouds for positive news

text = " ".join([x for x in air_data.tweet[air_data.sentiment == 'positive']])
wordcloud = WordCloud(background_color='white').generate(text)

plt.figure(figsize=(8,6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()

```



Jet Blue has the most Positive Tweets

```

# Negative Tweets

```






```
        return 0
    elif (st=='neutral'):
        return 2
    else:
        return 1

train['label'] = train['sentiment'].apply(making_label)
test['label'] = test['sentiment'].apply(making_label)

train_df = pd.DataFrame({
    'text': train['tweet'][:6000].replace(r'\n', '', regex=True),
    'label': train['label'][:6000]
})

eval_df = pd.DataFrame({
    'text': train['tweet'][-1000:].replace(r'\n', '', regex=True),
    'label': test['label'][-1000:]
})
bert_model.train_model(train_df)
```



```
sentiment
negative      9080
neutral       3057
positive      2290
Name: count, dtype: int64

(750, 0.5900795254607996)
```

### Multinomial Logistic Regression

```
# Multinomial Logistic Regression Analysis for tweets.

from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import CountVectorizer
import pandas as pd
import re, string
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import numpy as np
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
tweet_df =
pd.read_csv("C:/Users/coryg/OneDrive/Desktop/STAT_574_Data_Mining/hw4STAT574S25/USAirlinesTweets.csv")
tweet_df = tweet_df[['sentiment', 'tweet']].dropna()

def preprocessing_tweets(tweets):
    tweets = tweets.lower()
    tweets = re.sub(r"http\S+|www\S+|https\S+", '', tweets, flags=re.MULTILINE)
    tweets = re.sub(r'\@[\w]+\#', '', tweets)
    tweets = tweets.translate(str.maketrans('', '', string.punctuation))
    return tweets

tweet_df['cleaned_tweets'] = tweet_df['tweet'].apply(preprocessing_tweets)

# Text Vectorization to extract top 50 words from tweets

vectorization = CountVectorizer(max_features=50, stop_words='english')
```

```
X_mat = vectorization.fit_transform(tweet_df['cleaned_tweets'])
X_df = pd.DataFrame(X_mat.toarray(),
columns=vectorization.get_feature_names_out())
y = tweet_df['sentiment']
```

```
labeler = LabelEncoder()
y_encode = labeler.fit_transform(y)
```

```
# Split into 80% training and 20% testing sets and run the model.
```

```
X_train, X_test, y_train, y_test = train_test_split(X_df, y_encode,
test_size=0.2,
                                                    random_state=5720255)
```

```
multi_logistic = LogisticRegression(multi_class='multinomial', solver='lbfgs',
max_iter=900, random_state=987022)
multi_logistic.fit(X_train, y_train)
```

```
# Words most associated with Positive Tweets
```

```
coefs = pd.DataFrame(multi_logistic.coef_,
columns=vectorization.get_feature_names_out())
coefs.index = labeler.classes_
```

```
top_pos = coefs.loc['positive'].sort_values(ascending=False).head(10)
top_neg = coefs.loc['negative'].sort_values(ascending=False).head(10)
```

```
print("Positive tweets associated with the following words:")
print(top_pos)
print("\nNegative tweets associated with the following words:")
print(top_neg)
```

Positive tweets associated with the following words:

thank	2.010026
great	1.656956
thanks	1.540171
good	0.899002
airline	0.605125
guys	0.456910
got	0.424672
really	0.392620
flying	0.349597
home	0.331437

Name: positive, dtype: float64

Negative tweets associated with the following words:

hours	1.339060
hold	1.234786
delayed	1.215004
hour	1.163527
delay	0.990636
waiting	0.865747
bag	0.842824
cancelled	0.809081
phone	0.794711
dont	0.785270

Name: negative, dtype: float64

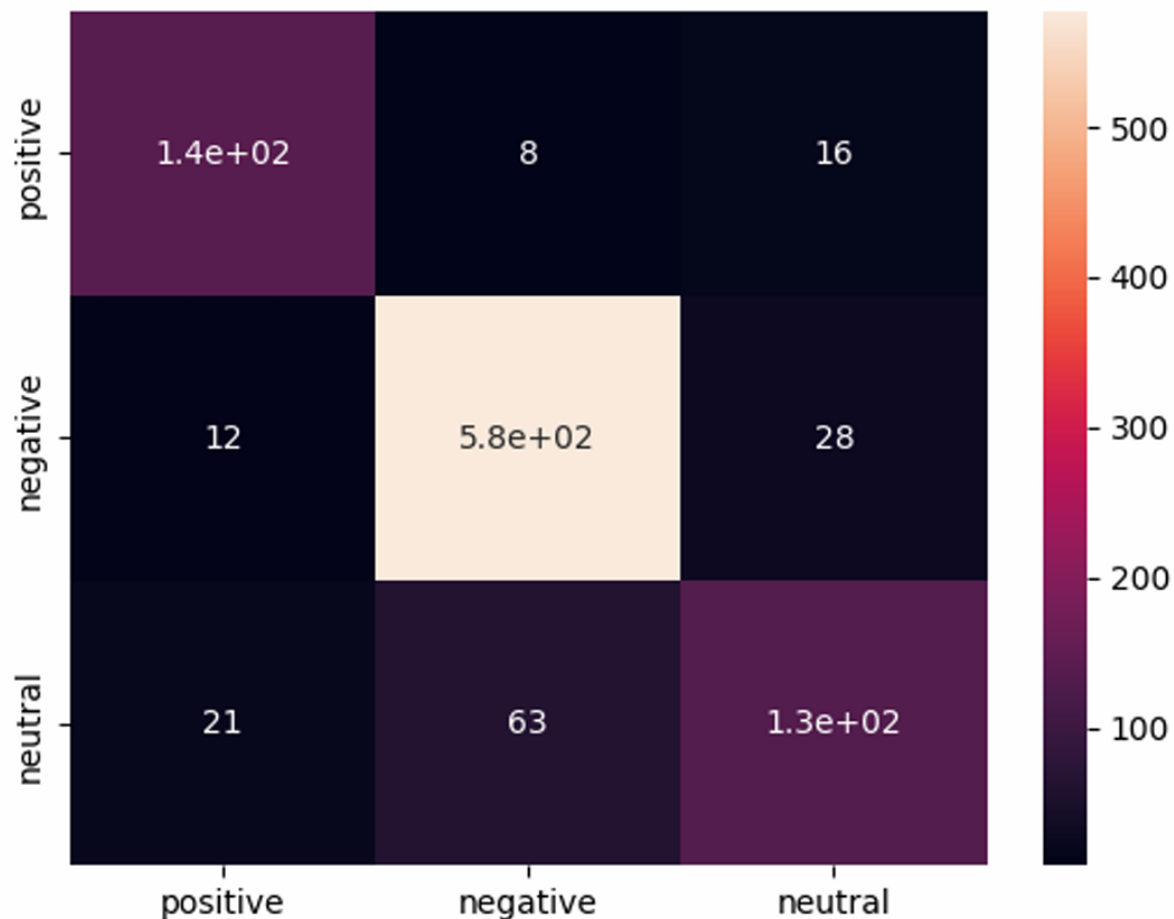
```
# Computing prediction accuracy on testing set.
```

```
y_pred = multi_logistic.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
print("accuracy: ", round(accuracy, 4))
print("confusion matrix: \n", confusion)
```

- accuracy: 0.68  
confusion matrix:  
[[1735 33 51]  
[ 543 41 43]  
[ 256 11 215]]

```
# Computing predicted sentiments for testing set.  
  
result, model_outputs, wrong_predictions = bert_model.eval_model(eval_df)  
lst = []  
for arr in model_outputs:  
    lst.append(np.argmax(arr))  
  
truth = eval_df['label'].tolist()  
predicted = lst  
  
# Displaying confusion matrix.  
  
conf_mat = confusion_matrix(truth, predicted)  
print(conf_mat)  
  
# Displaying heatmap for confusion matrix.  
  
df_cm = pd.DataFrame(conf_mat, ['positive', 'negative', 'neutral'], ['positive',  
'negative', 'neutral'])  
sns.heatmap(df_cm, annot=True)  
plt.show()
```

```
[[141 8 16]  
[ 12 577 28]  
[ 21 63 134]]
```



```
# Displaying performance metrics
```

```
print(sklearn.metrics.classification_report(truth, predicted,
target_names=['positive', 'negative', 'neutral']))
print(sklearn.metrics.accuracy_score(truth, predicted))
```

```
: '          precision    recall  f1-score   support\n\n   positive       0.85      0.83      0.84       165\n   negative       0.89      0.94      0.91       617\n   neutral       0.75      0.61      0.68       218\n\n accuracy\nmacro avg       0.85      0.85      0.85      1000\nweighted avg       0.85      0.85      0.85      1000'
```

0.852

```
# Using trained model to classify user-defined sentences.
```

```
def classify(statement):
    result = bert_model.predict([statement])
    pred_class = np.where(result[1][0] == np.amax(result[1][0]))
    pred_class = int(pred_class[0])
```

```
classify("Riding on a plane today.")
classify("The seats were cramped in my flight.")
classify("The crew were friendly and attentive.")
```

positive

```
def making_label(st):
    if (st == "Bad Flight"):
        return 0
    if (st == "Can't Tell"):
```



```

        return 1
    if (st == "Late Flight"):
        return 2
    if (st == "Customer Service Issue"):
        return 3
    if (st == "Flight Booking Problems"):
        return 4
    if (st == "Lost Luggage"):
        return 5
    if (st == "Flight Attendant Complaints"):
        return 6
    if (st == "Cancelled Flight"):
        return 7
    if (st == "Damaged Luggage"):
        return 8
    if (st == "longlines"):
        return 9

train['label'] = train['negativereason'].apply(making_label)
test['label'] = test['negativereason'].apply(making_label)

train_df = pd.DataFrame({
    'text': train['tweet'][:3500].replace(r'\n', ' ', regex=True),
    'label': train['label'][:3500]
})

eval_df = pd.DataFrame({
    'text': test['tweet'][-900:].replace(r'\n', ' ', regex=True),
    'label': test['label'][-900:]
})

neg_bert.train_model(train_df)

```

(438, 1.5106357800770023)

```

# Computing predicted sentiments for testing set.

result, model_outputs, wrong_predictions = neg_bert.eval_model(eval_df)

lst = []
for arr in model_outputs:
    lst.append(np.argmax(arr))

truth = eval_df['label'].tolist()

```

```

predicted = lst

# Computing predicted accuracy

print(accuracy_score(truth, predicted))

def classify(statement):
    result = neg_bert.predict([statement])
    pred_class = np.where(result[1][0] == np.amax(result[1][0]))
    pred_class = int(pred_class[0])
    sentiment_dict = {0: "Bad Flight", 1: "Can't Tell", 2: "Late Flight", 3:
"Customer Service Issue",
                     4: "Flight Booking Problems", 5: "Lost Luggage", 6: "Flight
Attendant Complaints",
                     7: "Cancelled Flight", 8: "Damaged Luggage", 9:
"longlines"}
    print(sentiment_dict[pred_class])
    return

classify('The staff were rude and disrespectful.')
classify('The flight was late.')
classify('My luggage got lost and was found damaged.')

```

0.63444444444444445

```

100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:09<00:00, 9.93s/it]
100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:00<00:00, 2.84it/s]
Customer Service Issue
100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:09<00:00, 9.43s/it]
100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:00<00:00, 2.90it/s]
Late Flight
100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:09<00:00, 9.41s/it]
100%|████████████████████████████████████████████████████████████████████████████████| 1/1 [00:00<00:00, 3.09it/s]
Lost Luggage

```